


Object-Oriented Programming Programming Project #1

Data Center

- A data center consists of multiple servers
- The servers are connected by switches in a local area network




Switches

- Each switch has multiple ports
- Receive and forward the packets from a port to another port




Traditional Routing Path

- Switches use OSPF (i.e., shortest path)
- Construct a shortest path tree rooted at each destination




Disadvantages of OSPF Routing Tables

- All paths are **fixed** → Not flexible
- **Periodical** update → Not real-time
- We need SDN-enabled switches

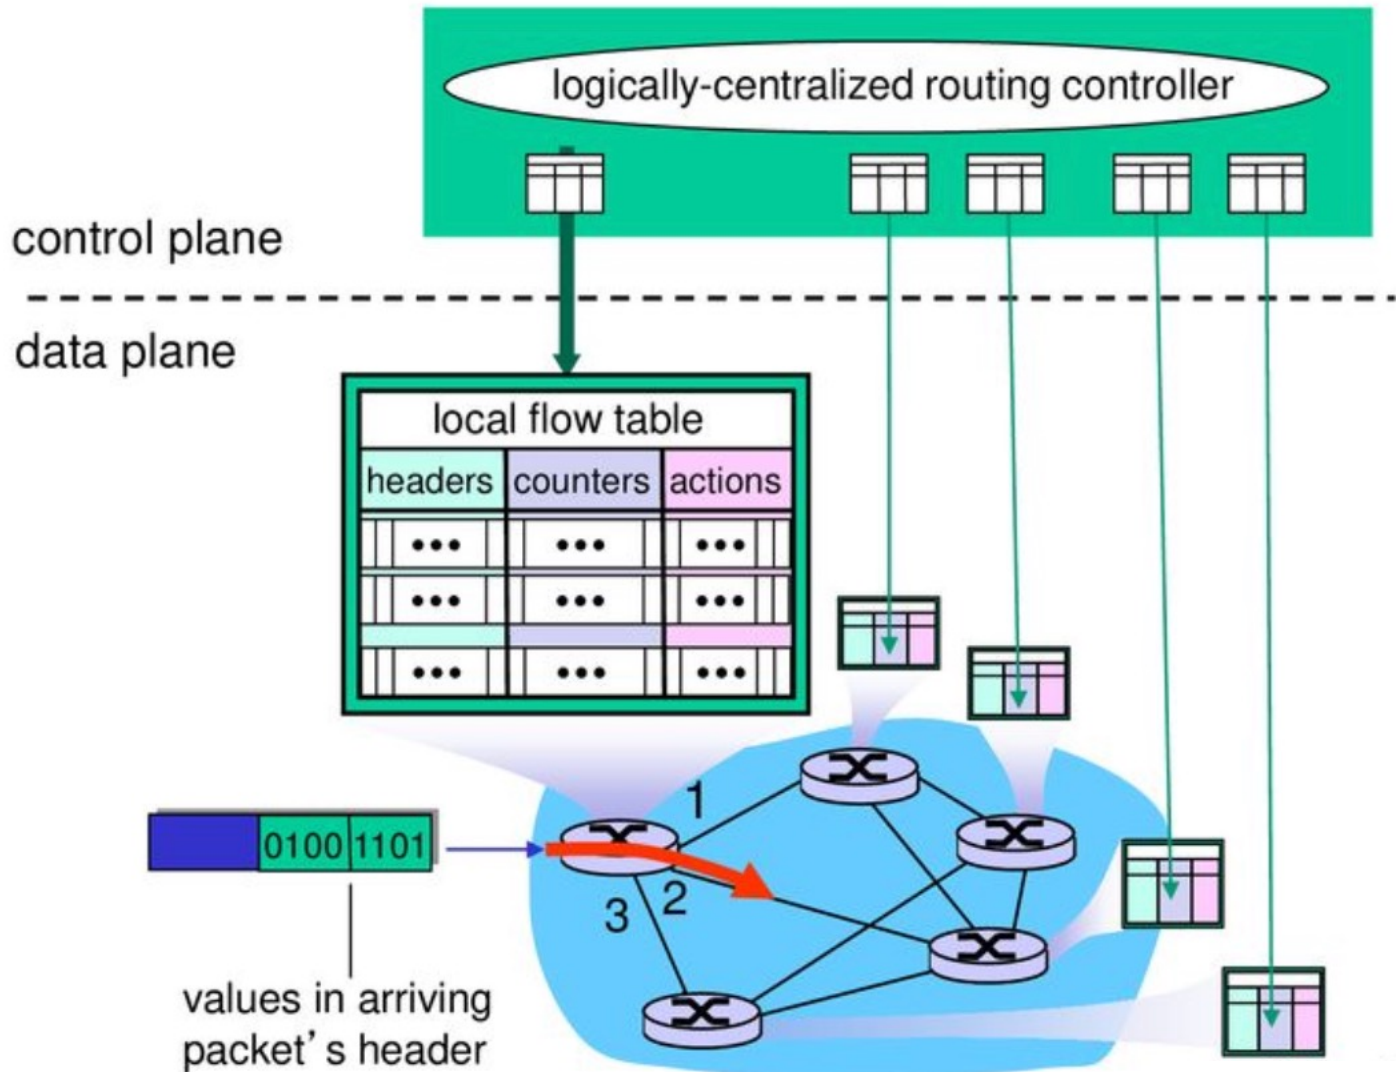


SDN-enabled Switches

- A centralized controller is introduced – software-defined networking (**SDN**)




Installing Rules in the SDN-enabled Switches




Incrementally Deployed SDN

- Non-SDN switches: uncontrollable OSPF paths
- **SDN switches**: controllable paths



OSPF Routing Information


- Given: a graph with links and destinations
- Output: shortest paths towards all destinations
- Then, store the information in each node's table



OSPF Routing Table

- Key: each **destination**
- Value: the **next node** (i.e., the output port)
- **Node 1's table** (it uses OSPF)

Destination	Next Node
0	2



Routing Flows with OSPF Routing Tables


- Given flows:

- Flow 1 \rightarrow 0 with flow size 3
- Flow 2 \rightarrow 0 with flow size 4
- Flow 3 \rightarrow 0 with flow size 5

- Links' loads:

- Link (1, 2)'s load = 3
- Link (3, 2)'s load = 5
- Link (2, 10)'s load = 12

- Max link load
 $= \max\{3, 5, 12\}$
 $= 12$




SDN-enabled Routing Table

- Key: each **destination**
- Value: the **next nodes** (i.e., the output ports)
- **Node 2's table** (it is SDN-enabled)

Destination	Next Node	Portion
0	10, 11	50%, 50%

The controller can define the portions



Routing Flows with Hybrid Routing Tables


- Given flows:

- Flow 1 \rightarrow 0 with flow size 3
- Flow 2 \rightarrow 0 with flow size 4
- Flow 3 \rightarrow 0 with flow size 5

- Links' loads:

- Link (1, 2)'s load = 3
- Link (3, 2)'s load = 5
- Link (2, 10)'s load = 6
- Link (2, 11)'s load = 6

- Max link load
= $\max \{3, 5, 6, 6\}$
= 6




Requirements

- Every node only knows its neighbors
- Define your own **node class** and use **public** (i.e., just use class instead of struct)
- Each node has an unsigned int ID
- Use a **vector<unsigned int>** to store the **neighbors' IDs** in each node
- Use a **map<unsigned int, vector<pair<unsigned int, double> > >** to store **each entry** in the table (i.e., each entry in the table has destination ID, <next nodes' IDs, portions>)

Note

- Serious **congestion problems** happen if cycles exist
- Avoid **cycles** in the routing paths for each destination



Programming Project #1:

Routing Table in an Incrementally Deployed SDN

- Input:

- # nodes, #SDN nodes, # destinations, # links, and #pairs
- SDN nodes (ID)
- Destinations (ID)
- Links between nodes
- Traffic matrix (flow size for each pair)

- Procedure:

- Compute shortest paths to each destination
- Compute next hops and portions for SDN-enabled nodes

- Output:

- Each node's routing table

The Competition

- The grade is inversely proportional to **the max link load**
- **Basic: 60 (deadline)**
 - Every node's packet can be sent to the destination with no cycle
- **Being a coding assistant (superb deadline)**
 - +10
- **Performance ranking (decided after the deadline)**
 - [0%, 30%) (bottom): +0
 - [30%, 50%): + 5
 - [50%, 75%): + 10
 - [75%, 85%): + 15
 - [85%, 90%): + 20
 - [90%, 95%): + 25
 - [95%, 100%] (top): + 30

The Competition

- Note that you **cannot** use brute-force algorithm
- Note that your code must be deterministic (no **randomization**)



Input Sample:

use cin

Format:

#Nodes #SDN_Nodes #Dsts #Links #Pairs

SDN_NodeID_List

DstID_List


...

LinkID Node1 Node2

...

FlowID Src Dst FlowSize

...



15 3 1 28 3

2 9 14

0

0 0 5

1 0 10

2 0 12

3 0 14

4 1 2

5 1 3

6 1 4

7 2 3

8 2 10

9 2 11

10 3 4

11 3 6

12 3 7

13 4 8

14 4 9

15 5 6

16 5 7

17 5 9

18 5 10

19 5 14

20 6 7

21 6 10

22 7 9

23 8 9

24 9 13

25 10 12

26 11 12

27 13 14

0 1 0 3

1 2 0 4

2 3 0 5

Output Sample (not optimal):

use cout

Format:

NodeID

DstID NextID


...

e.g.,

0
0 0
1
0 2
2
0 10 50% 11 50%
3
0 2
4
0 9
5
0 0
6
0 5

Its own ID

7
0 5
8
0 9
9
0 5 60% 7 0% 13 40%
10
0 0
11
0 12
12
0 0
13
0 14
14
0 0 70% 5 30%



Note

- Superb deadline: 3/16 Thu
- Deadline: 3/23 Thu
- Pass the test of our [online judge](#) platform
- Submit your code to [E-course2](#)
 - The file name should be ``OOP_HW1_studentID.cpp``
- Demonstrate your code [remotely](#) with TA
- **C++ Source code (only C++; compiled with g++)**
 - Include C++ library only (i.e., no stdio, no stdlib, ...)
 - Please use new and delete instead of malloc and free
- Show a good programming style