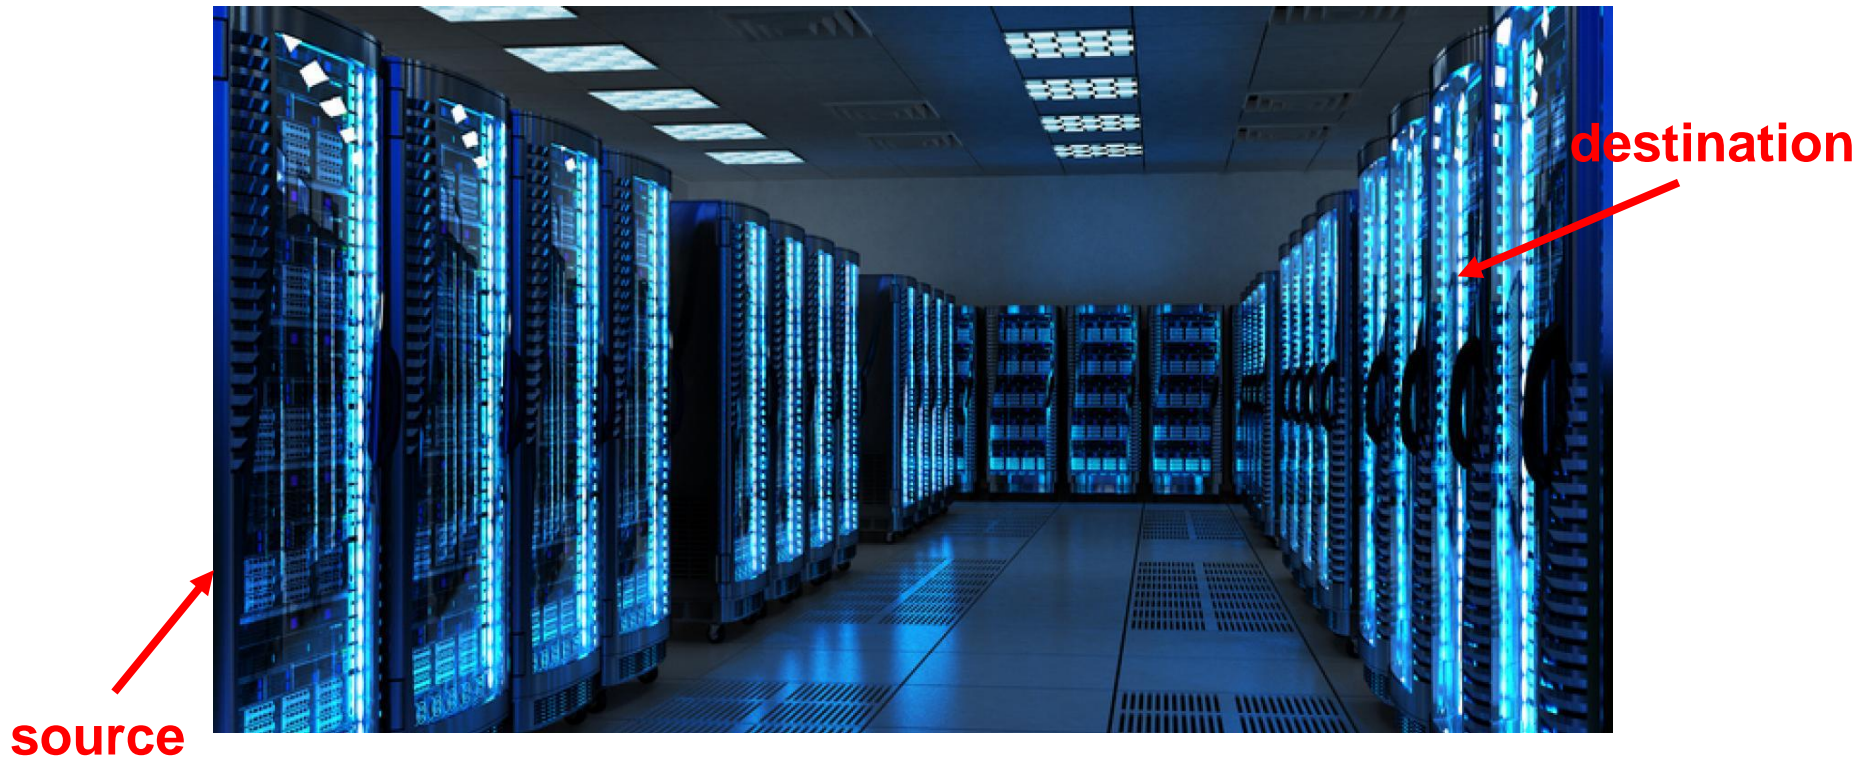


Object-Oriented Programming Programming Project #4

Data Center

- A data center consists of multiple servers
- The servers are connected by switches in a local area network




Switches

- Each switch has multiple ports
- Receive and forward the packets from a port to another port




Traditional Routing Path

- Switches use OSPF (i.e., shortest path)
- Construct a shortest path tree rooted at each destination



OSPF Routing Information (HW2)


- Given: a graph with links and destinations
- Output: shortest paths towards all destinations
- Then, store the information in each node's table



OSPF Routing Table


- Key: each **destination**
- Value: the **next node** (i.e., the output port)
- **Node 1's table** (it uses OSPF)

Destination	Next Node
0	2




Note – Rule for Selecting the Next Hop (HW2)

- Select the node with **a smaller counter** (i.e., closer to the destination)
- Select the node with **a smaller ID** as the next hop **if there is a tie** (i.e., multiple candidates)




Note – Rule for Relaying TRA_ctrl_packet (HW2)

- Relay packets with **a counter smaller than all my currently received counters**
- Relay packets with **a counter equal to my next hop's counter but with a preID smaller than my current next hop**




Splitting Flows with SDN switches (New)

- Given **a set of SDN switches**
- Goal: leverage the SDN switches to make the maximum link load **smaller than** that with no SDN switch




Max link load = 12



Max link load = 6

Note – No Cycle for Routing (HW1)

- Serious **congestion problems** happen if cycles exist
- Avoid **cycles** in the routing paths for each destination



Programming Project #4:

Decrease Max Link Load with SDN switches

- Input:
 - # nodes, # destinations, # links, and #pairs
 - SDN nodes (ID)
 - Destinations (ID)
 - Links between nodes
 - Traffic matrix (flow size for each pair)
- Procedure:
 - Compute shortest paths to each destination in a distributed manner
 - Invoke SDN controller to compute next hops and portions for SDN-enabled nodes (a must)
 - Route packets with different size toward destinations
- Output:
 - Each node's routing table
 - Packet exchange information will be logged automatically

Note – Create TRA Nodes and Links (HW2)

- Create traditional switches (i.e., `class TRA_switch`)
- Each traditional switch has an unsigned int `ID`
 - `node::node_generator::generate("TRA_switch", id);`
- Every node only `knows its neighbors`
- `Add the neighbors` for each traditional switch
 - `node::id_to_node(0)->add_phy_neighbor(1);`
 - `node::id_to_node(1)->add_phy_neighbor(0);`
 - We use `simple_link` with a fixed latency (i.e., 10)
- Write a `map<unsigned int, unsigned int>` to store `each entry` in each switch's table (i.e., each entry in the table has `<a destination ID, a next node ID>`) in class `TRA_switch`
 - Copy and modify partial code in HW2

Note – Create SDN Nodes and Links (HW3)

- Define class **SDN_switch** and **Create SDN_switch**
 - Derived from **class node (Inheritance)**
- Each SDN switch has an unsigned int **ID**
 - `node::node_generator::generate("SDN_switch",id);`
- Every node only **knows its neighbors**
- **Add the neighbors** for each SDN switch
 - `node::id_to_node(0)->add_phy_neighbor(1);`
 - `node::id_to_node(1)->add_phy_neighbor(0);`
 - We use `simple_link` with a fixed latency (i.e., 10)
- Write a **`map<unsigned int, vector<pair<unsigned int, double>>>`** to store **each entry** in each switch's table (i.e., each entry in the table has destination ID, <next nodes' IDs, portions>) in class **SDN_switch**
 - Copy and modify the routing table code in HW3

Note – Define and Create SDN Controller (HW3)

- Define new **class SDN_controller**
 - Derived from **class node (Inheritance)**
- After creating the switches, create an **SDN controller**
 - con_id is the controller ID (after all switches' ID)
 - node_generator::generate("SDN_controller", con_id);
- **Connect each SDN switch** to the controller
 - node::id_to_node(switch_id)->add_phy_neighbor(con_id);
 - node::id_to_node(con_id)->add_phy_neighbor(switch_id);

Note – Generate Traditional Ctrl Packets (HW2)

- **Generate traditional ctrl packets**


- void TRA_ctrl_packet_event(unsigned int src, unsigned int t = event::getCurTime(), string msg = "default")
- The function is used to initialize the distributed BFS; that is, a TRA_ctrl_packet will be **generated for a source (src) with a counter 0**
- You have to implement recv_handler() in TRA_switch to forward the ctrl packet to the neighboring node again; that is, **every node receiving the packet should increase the counter and broadcast the packet to its neighboring nodes** to update the rule in every node's table to build the path from every node to the source
- The source (src) will receive the TRA_ctrl_packet first (since it's src)

Note – Receive and Send Packets (1/2) (HW2/3)

- Define the rules to **handle the received packet** in class TRA_switich/SDN_switch's member function recv_handler
 - void TRA_switch::recv_handler (packet *p) (or SDN_switch's)
 - **Don't use node::id_to_node(id) in recv_handler**
- Get the **current switch's ID and its neighbor**
 - Use getNodeID() in recv_handler
 - Use getPhyNeighbors().find(n_id) to check whether the node with n_id is a neighbor
 - Use const map<unsigned int,bool> &nblast =getPhyNeighbors() and for (map<unsigned int,bool>::const_iterator it = nblast.begin(); it != nblast.end(); it ++) to get all neighbors
- Use **send_handler(packet *p)** to send the packet *p
- Check the packet type
 - if (p->type() == "TRA_data_packet")
 - if (p->type() == "TRA_ctrl_packet")
 - if (p->type() == "SDN_ctrl_packet")

Note – Receive and Send Packets (2/2) (HW2)

- Decode: **Cast the packet, payload**, to the right type
 - `TRA_data_packet *p2 = dynamic_cast<TRA_data_packet*> (p)`
 - `TRA_ctrl_packet *p3 = dynamic_cast<TRA_ctrl_packet*> (p)`
 - `TRA_ctrl_payload *l3 = dynamic_cast<TRA_ctrl_payload*> (p3->getPayload());`
 - ...
- **Before sending** a packet to the next hop
 - Use `setPreID(id)` to change the preID to the current node's ID
 - Use `setNexID(id)` to change the nexID to the next hop node's ID
 - Please check all the columns in the header



Note – Generate SDN Ctrl Packets (HW3)

- **Generate SDN ctrl packets**

- `void SDN_ctrl_packet_event(unsigned int con_id, unsigned int id, unsigned int mat, unsigned int act, double per, unsigned int t = event::getCurTime(), string msg = "default")`
- A packet will be **generated for the controller** and sent to the node (id) **at time t (optional)** to update a specific rule in the node's table
- **mat**: the destination in the node's routing table
act: the next hop in the node's routing table
per: the portion of flow sent to the next node
- The node can use **getMatID(), getActID(), and getPer()** of `SDN_ctrl_payload` to get mat, act, and per from the `SDN_ctrl_packet`
- The controller will receive the `SDN_ctrl_packet` first (since the controller is src)

Note – Invoke Computing Portions (New)

- **Invoke SDN controller to compute portions**

- `void SDN_invoke_packet_event(unsigned int con_id, vector<vector<double>> traffic_matrix, unsigned int t = event::getCurTime(), string msg = "default")`
- A packet will be generated for the controller at time t (optional) to notify the controller to compute the portions for each SDN switch based on the traffic matrix (traffic_matrix)
- The controller will receive the SDN_invoke_packet first (since the controller is src and also dst)
- The controller can know that the SDN_invoke_packet is used to notify itself to compute portions
- Move the code for computing routing tables in SDN switches into the recv_handler function of the SDN controller
- The function `recv_handler` in SDN controller is allowed to use `node::id_to_node(id)`

Note – Generate Data and Ctrl Packets (New)

- Generate data packets

- void **data_packet_event**(unsigned int src, unsigned int dst, double size, unsigned int t = 0, string msg="default")
- A TRA_data_packet with **the flow size (size)** will be **generated for a source (src)** and **sent to a destination (dst) at time t**
- The source (src) will receive the TRA_data_packet first (since it's src)

Inheritance

packet

TRA_ctrl_packet

TRA_data_packet

SDN_ctrl_packet

SDN_invoke_packet

packet class
Methods
discard
getLivePacketNum
~packet

SDN_ctrl_packet class
Methods
type
~SDN_ctrl_packet

TRA_ctrl_packet class
Methods
type
~TRA_ctrl_packet

TRA_data_packet class
Methods
type
~TRA_data_packet

packet_generator class
Methods
generate
print
replicate
~packet_generator

SDN_ctrl_packet_generator class
Methods
type
~SDN_ctrl_packet_generator

TRA_ctrl_packet_generator class
Methods
type
~TRA_ctrl_packet_generator

TRA_data_packet_generator class
Methods
type
~TRA_data_packet_generator

Inheritance









header

TRA_ctrl_header

TRA_data_header

SDN_ctrl_header

SDN_invoke_header



Inheritance

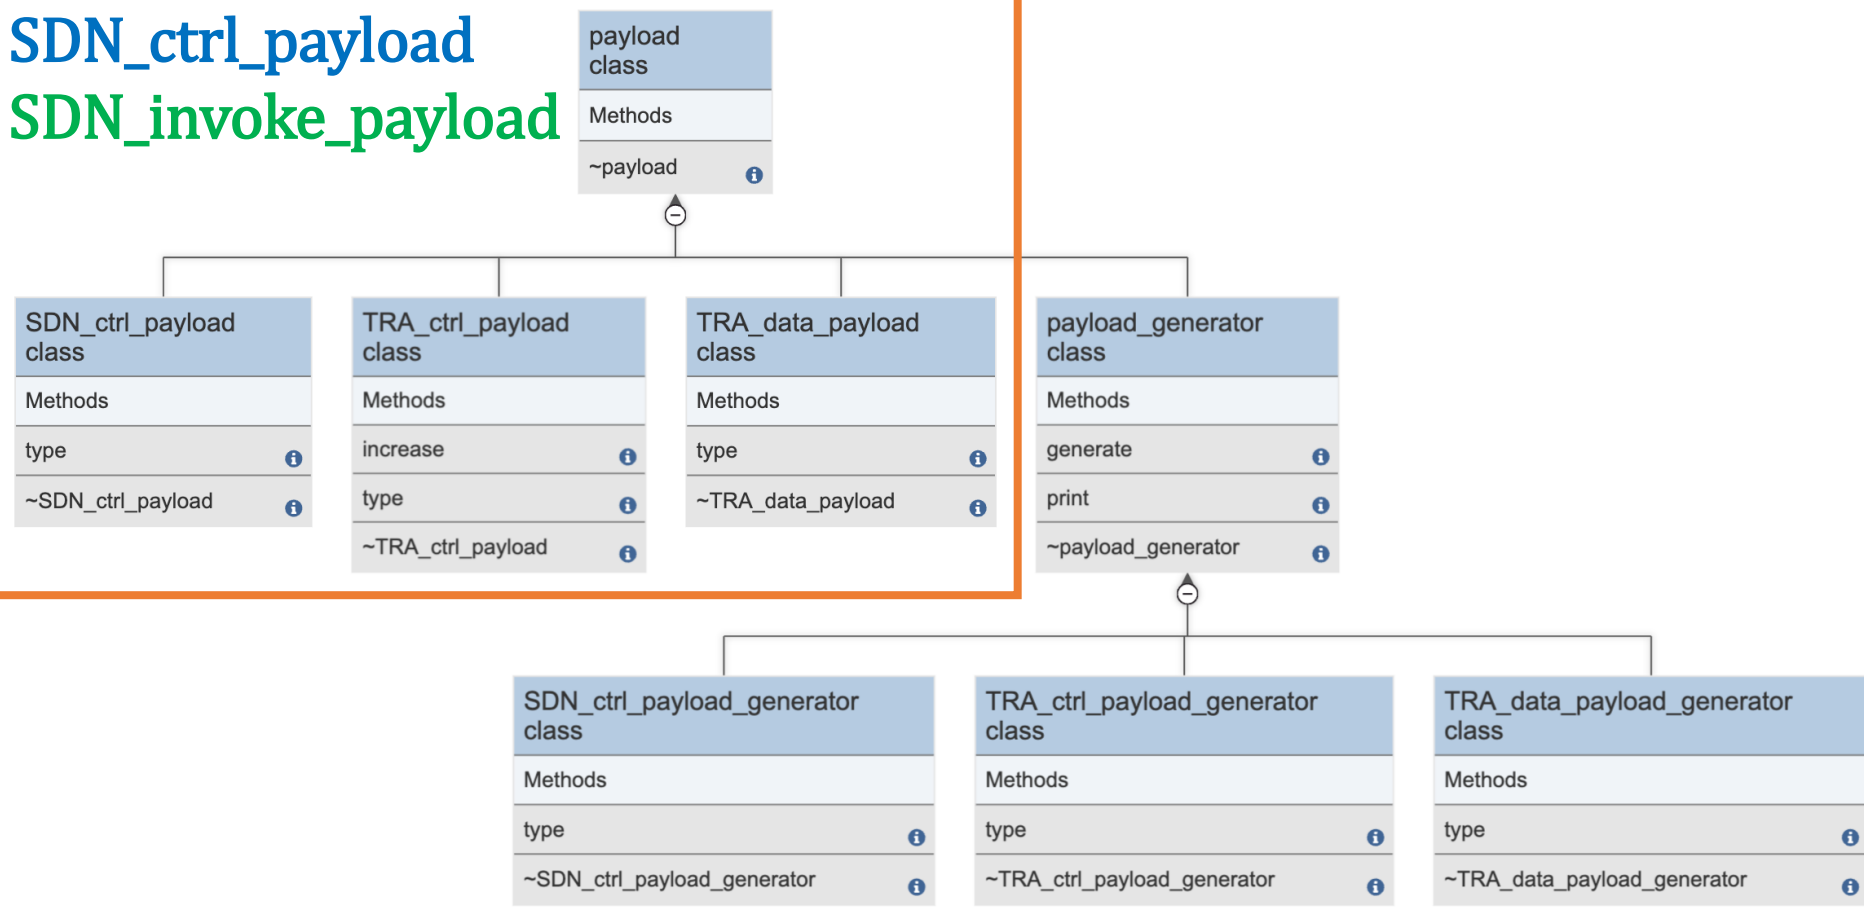
payload

TRA_ctrl_payload

TRA_data_payload

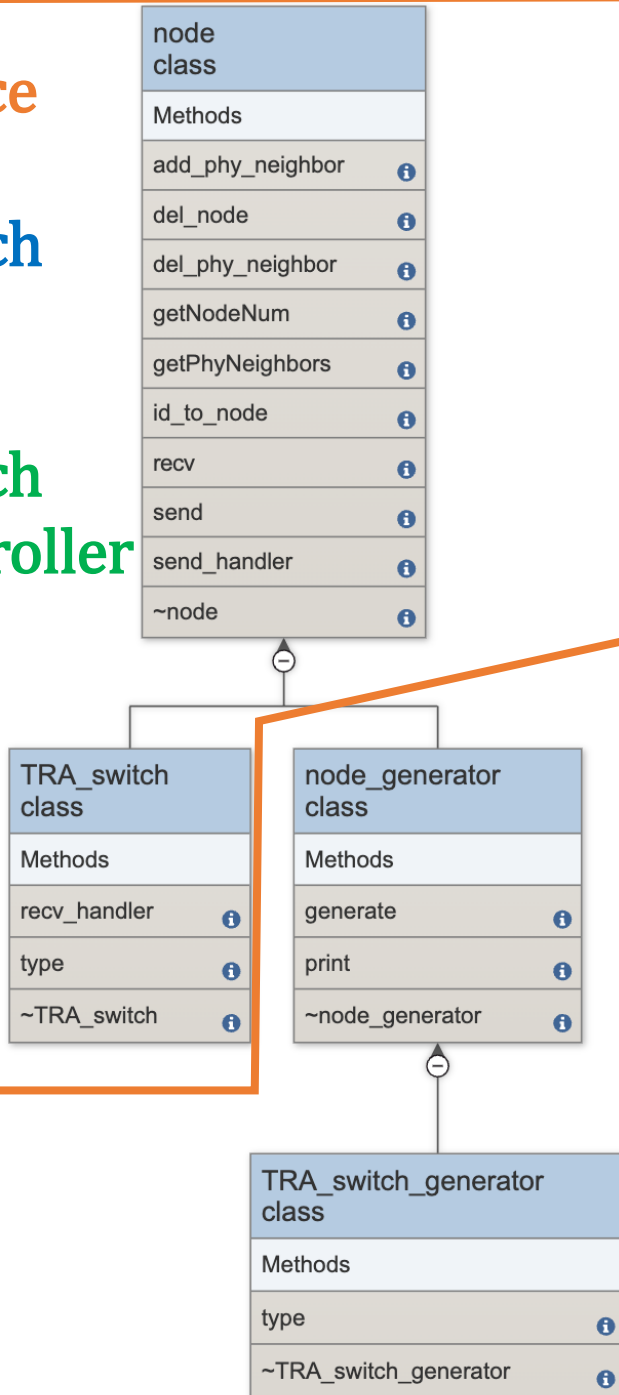
SDN_ctrl_payload

SDN_invoke_payload

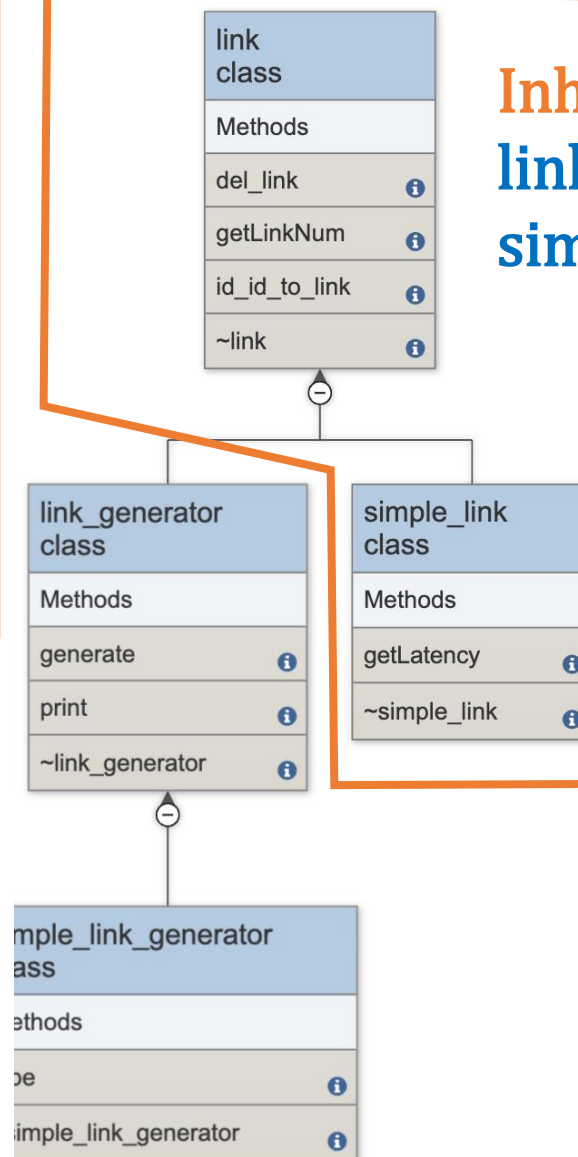


Inheritance
node
TRA_switch

Required:
SDN_switch
SDN_controller



Inheritance
link
simple_link



Input Sample:

use cin

Format:

#Nodes #SDN_Nodes #Dsts #Links #Pairs SimTime InvokeTime

SDN_NodeID_List

DstID_List

NodeID BroadcastTime


...

LinkID Node1 Node2

...

FlowID Src Dst FlowSize StartTime

...



```
0 0 5
1 0 10
2 0 12
3 0 14
4 1 2
5 1 3
6 1 4
7 2 3
8 2 10
9 2 11
10 3 4
11 3 6
12 3 7
13 4 8
14 4 9
15 5 6
16 5 7
17 5 9
18 5 10
19 5 14
20 6 7
21 6 10
22 7 9
23 8 9
24 9 13
25 10 12
26 11 12
27 13 14
0 1 0 3 200
1 2 0 4 220
2 3 0 5 250
```

Output Sample:

use cout

You have to print the routing table for each node after event::start_simulate();

The way to print the routing table is the same as that in HW1/HW2/HW3

The remaining output will be automatically generated 😊

Note that the output could be different in different computers

use cout

Format:

UpgradedNodeIDList

DstID	NextID	Portion	NextID	Portion...
-------	--------	---------	--------	------------

The graph consists of 15 nodes labeled 0 through 14. Nodes 2, 0, and 9 are highlighted in blue. The connections are as follows:

- Thick Yellow Arrows (Primary Flow):**
 - 1 → 2
 - 2 → 11
 - 2 → 10
 - 11 → 12
 - 12 → 0
 - 0 → 14
 - 14 → 13
 - 13 → 9
 - 9 → 5
 - 5 → 7
 - 7 → 6
 - 6 → 3
 - 3 → 4
 - 4 → 8
 - 8 → 9
- Thin Black Arrows (Other Connections):**
 - 1 → 4
 - 3 → 7
 - 4 → 3
 - 5 → 10
 - 10 → 6
 - 9 → 14
 - 13 → 5
 - 14 → 0

The graph is partially truncated on the left, indicated by an ellipsis (...).

2 9 14

0 0

0 2

0 10 50% 11 50%

0 2

0 9

0 0

0 5

0 5 60% 7 0% 13 40%

0 0

0 12
12

13

14
2

0 0 70% 5 50%

Its own ID

Note

- Superb deadline: 6/18 Sun (you have more than 2 weeks)
- Deadline: 6/21 Wed (you have about 3 weeks)
- Pass the test of our [online judge](#) platform
- Submit your code to [E-course2](#)
- Demonstrate your code [remotely](#) with TA (before 6/18 or 6/21)
- **C++ Source code (only C++; compiled with g++)**
 - Please use C++ library (i.e., no stdio, no stdlib)
 - Please use new and delete instead of malloc and free
- Show a good programming style

